

Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [1]: 1 # Run some setup code for this notebook.
2
3 from __future__ import print_function
4 import random
5 import numpy as np
6
7 from cs231n.data_utils import load_CIFAR10
8 import matplotlib.pyplot as plt
9
10
11 # This is a bit of magic to make matplotlib figures appear inline in the
12 # notebook rather than in a new window.
13 %matplotlib inline
14 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
15 plt.rcParams['image.interpolation'] = 'nearest'
16 plt.rcParams['image.cmap'] = 'gray'
17
18 # Some more magic so that the notebook will reload external python modules;
19 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
20 %load_ext autoreload
21 %autoreload 2
```

CIFAR-10 Data Loading and Preprocessing

In [2]:

```
1 # Load the raw CIFAR-10 data.
2 cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
3
4 # Cleaning up variables to prevent loading data multiple times (which may cause memory issue)
5 try:
6     del X_train, y_train
7     del X_test, y_test
8     print('Clear previously loaded data.')
9 except:
10     pass
11
12 X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
13
14 # As a sanity check, we print out the size of the training and test data.
15 print('Training data shape: ', X_train.shape)
16 print('Training labels shape: ', y_train.shape)
17 print('Test data shape: ', X_test.shape)
18 print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

In [3]:

```
1 # Visualize some examples from the dataset.
2 # We show a few examples of training images from each class.
3 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'] #类别列表
4 num_classes = len(classes) #类别数目
5 samples_per_class = 7 # 每个类别采样个数
6 for y, cls in enumerate(classes): ## 对列表的元素位置和元素进行循环, y表示元素位置[0,num_class), cls元素本身'plane'等
7     idxs = np.flatnonzero(y_train == y) #找出标签中y类的位置
8     idxs = np.random.choice(idxs, samples_per_class, replace=False) #从中选出我们所需的7个样本
9     for i, idx in enumerate(idxs): #对所选的样本的位置和样本所对应的图片在训练集中的位置进行循环
10         plt_idx = i * num_classes + y + 1 # 在子图中所占位置的计算
11         plt.subplot(samples_per_class, num_classes, plt_idx) # 说明要画的子图的编号
12         plt.imshow(X_train[idx].astype('uint8')) # 画图
13         plt.axis('off')
14         if i == 0:
15             plt.title(cls) # 写上标题, 也就是类别名
16 plt.show() # 显示
```



```
In [4]: 1 # Split the data into train, val, and test sets. In addition we will
2 # create a small development set as a subset of the training data;
3 # we can use this for development so our code runs faster.
4 num_training = 49000
5 num_validation = 1000
6 num_test = 1000
7 num_dev = 500
8
9 # Our validation set will be num_validation points from the original
10 # training set.
11 mask = range(num_training, num_training + num_validation) ##[49000,50000)
12 X_val = X_train[mask]
13 y_val = y_train[mask]
14
15 # Our training set will be the first num_train points from the original
16 # training set.
17 mask = range(num_training) ####[0,49000)
18 X_train = X_train[mask]
19 y_train = y_train[mask]
20
21 # We will also make a development set, which is a small subset of
22 # the training set.
23 mask = np.random.choice(num_training, num_dev, replace=False) #shape500
24 X_dev = X_train[mask]
25 y_dev = y_train[mask]
26
27 # We use the first num_test points of the original test set as our
28 # test set.
29 mask = range(num_test)
30 X_test = X_test[mask]
31 y_test = y_test[mask]
32
33 print('Train data shape: ', X_train.shape)
34 print('Train labels shape: ', y_train.shape)
35 print('Validation data shape: ', X_val.shape)
36 print('Validation labels shape: ', y_val.shape)
37 print('Test data shape: ', X_test.shape)
38 print('Test labels shape: ', y_test.shape)
```

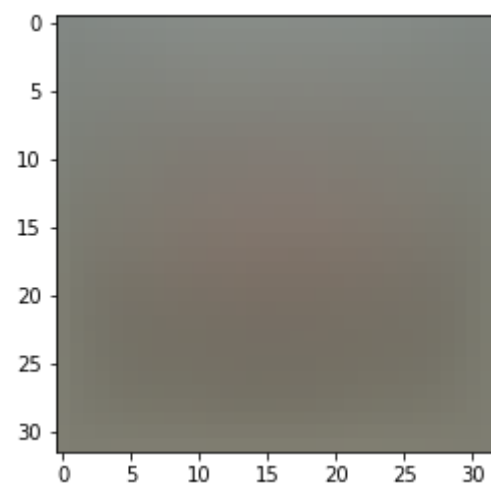
```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```

```
In [5]: 1 # Preprocessing: reshape the image data into rows
2 X_train = np.reshape(X_train, (X_train.shape[0], -1)) # newshape == (X_train.shape[0], -1) == (x, m*n/x) == (49000,3072)
3 X_val = np.reshape(X_val, (X_val.shape[0], -1))
4 X_test = np.reshape(X_test, (X_test.shape[0], -1))
5 X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))
6
7 # As a sanity check, print out the shapes of the data
8 print('Training data shape: ', X_train.shape)
9 print('Validation data shape: ', X_val.shape)
10 print('Test data shape: ', X_test.shape)
11 print('dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

```
In [6]: 1 # Preprocessing: subtract the mean image
2 # first: compute the image mean based on the training data
3 mean_image = np.mean(X_train, axis=0)
4 print(mean_image[:10]) # print a few of the elements
5 plt.figure(figsize=(4,4))
6 plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean image
7 plt.show()
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
 131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



```
In [7]: 1 # second: subtract the mean image from train and test data
2 X_train -= mean_image
3 X_val -= mean_image
4 X_test -= mean_image
5 X_dev -= mean_image
```

```
In [8]: 1 # third: append the bias dimension of ones (i.e. bias trick) so that our SVM
2 # only has to worry about optimizing a single weight matrix W.
3 X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))]) #Bias trick
4 X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
5 X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
6 X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])
7
8 print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)

(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)
```

SVM Classifier

Your code for this section will all be written inside **cs231n/classifiers/linear_svm.py**.

As you can see, we have prefilled the function `compute_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
In [9]: 1 # Evaluate the naive implementation of the loss we provided for you:
2 from cs231n.classifiers.linear_svm import svm_loss_naive
3 import time
4
5 # generate a random SVM weight matrix of small numbers
6 W = np.random.randn(3073, 10) * 0.0001 ## 3073是数据集的特征数目3072再加上1,w中多出来的这一列对应的就是偏差值b
7
8 #W.shape == (D,C) == (3073, 10) X_dev.shape == (N,D) == (500,3073) y_dev.shape == (N,) == (500,)
9 loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
10 print('loss: %f' % (loss, ))

loss: 8.943680
```

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

```
In [10]: 1 # Once you've implemented the gradient, recompute it with the code below
2 # and gradient check it with the function we provided for you
3
4 # Compute the loss and its gradient at W.
5 loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)
6
7 # Numerically compute the gradient along several randomly chosen dimensions, and
8 # compare them with your analytically computed gradient. The numbers should match
9 # almost exactly along all dimensions.
10 from cs231n.gradient_check import grad_check_sparse
11 f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
12 grad_numerical = grad_check_sparse(f, W, grad)
13
14 # do the gradient check once again with regularization turned on
15 # you didn't forget the regularization gradient did you?
16 loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
17 f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
18 grad_numerical = grad_check_sparse(f, W, grad)
```

```
numerical: 35.769934 analytic: 35.769934, relative error: 7.428433e-12
numerical: 10.757289 analytic: 10.757289, relative error: 3.039814e-11
numerical: 23.510778 analytic: 23.510778, relative error: 1.080096e-11
numerical: -0.769543 analytic: -0.769543, relative error: 8.175079e-11
numerical: 1.380252 analytic: 1.380252, relative error: 1.623525e-12
numerical: 2.838039 analytic: 2.838039, relative error: 2.467822e-11
numerical: 21.331356 analytic: 21.331356, relative error: 5.570144e-12
numerical: 6.749857 analytic: 6.749857, relative error: 9.462458e-12
numerical: 0.072874 analytic: 0.072874, relative error: 1.672515e-09
numerical: -14.212107 analytic: -14.212107, relative error: 4.358741e-12
numerical: -27.798089 analytic: -27.796582, relative error: 2.711249e-05
numerical: -36.294893 analytic: -36.298383, relative error: 4.807381e-05
numerical: 5.247237 analytic: 5.247207, relative error: 2.789185e-06
numerical: -8.238703 analytic: -8.239616, relative error: 5.542042e-05
numerical: -10.086545 analytic: -10.084483, relative error: 1.022036e-04
numerical: 16.459586 analytic: 16.459600, relative error: 4.199611e-07
numerical: -32.731761 analytic: -32.723241, relative error: 1.301718e-04
numerical: 0.874677 analytic: 0.880805, relative error: 3.490755e-03
numerical: 25.547095 analytic: 25.541069, relative error: 1.179465e-04
numerical: -51.988328 analytic: -51.988321, relative error: 6.908752e-08
```

Inline Question 1:

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect of the frequency of this happening? *Hint: the SVM loss function is not strictly speaking differentiable*

Your Answer: *fill this in.*

```
In [11]: 1 # Next implement the function svm_loss_vectorized; for now only compute the loss;
2 # we will implement the gradient in a moment.
3 tic = time.time()
4 loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
5 toc = time.time()
6 print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))
7
8 from cs231n.classifiers.linear_svm import svm_loss_vectorized
9 tic = time.time()
10 loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
11 toc = time.time()
12 print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))
13
14 # The losses should match but your vectorized implementation should be much faster.
15 print('difference: %f' % (loss_naive - loss_vectorized))
```

```
Naive loss: 8.943680e+00 computed in 0.275538s
Vectorized loss: 8.943680e+00 computed in 0.017083s
difference: 0.000000
```

```
In [12]: 1 # Complete the implementation of svm_loss_vectorized, and compute the gradient
2 # of the loss function in a vectorized way.
3
4 # The naive implementation and the vectorized implementation should match, but
5 # the vectorized version should still be much faster.
6 tic = time.time()
7 _, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
8 toc = time.time()
9 print('Naive loss and gradient: computed in %fs' % (toc - tic))
10
11 tic = time.time()
12 _, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
13 toc = time.time()
14 print('Vectorized loss and gradient: computed in %fs' % (toc - tic))
15
16 # The loss is a single number, so it is easy to compare the values computed
17 # by the two implementations. The gradient on the other hand is a matrix, so
18 # we use the Frobenius norm to compare them.
19 difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
20 print('difference: %f' % difference)
```

```
Naive loss and gradient: computed in 0.454984s
Vectorized loss and gradient: computed in 0.006603s
difference: 0.000000
```

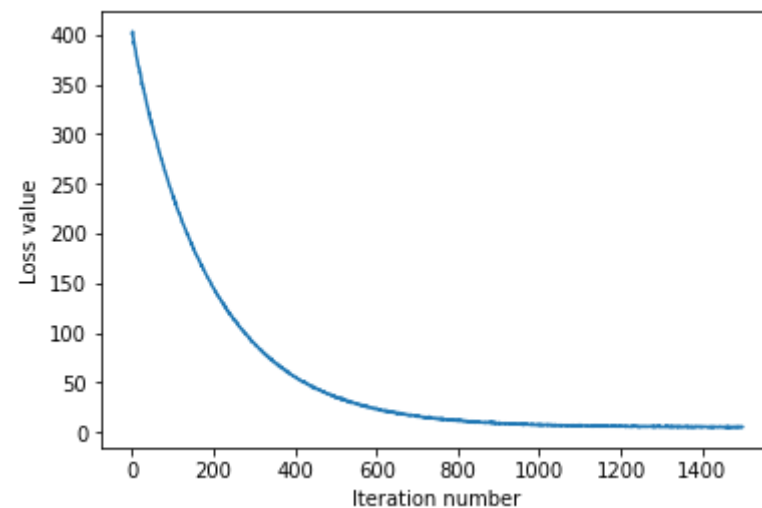
Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss.


```
In [13]: 1 # In the file linear_classifier.py, implement SGD in the function
2 # LinearClassifier.train() and then run it with the code below.
3 from cs231n.classifiers import LinearSVM
4 svm = LinearSVM()
5 tic = time.time()
6 loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
7                       num_iters=1500, verbose=True)
8 toc = time.time()
9 print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 401.179673
iteration 100 / 1500: loss 236.957475
iteration 200 / 1500: loss 145.353562
iteration 300 / 1500: loss 89.138484
iteration 400 / 1500: loss 55.495465
iteration 500 / 1500: loss 35.054152
iteration 600 / 1500: loss 24.282576
iteration 700 / 1500: loss 16.155006
iteration 800 / 1500: loss 11.949770
iteration 900 / 1500: loss 9.582905
iteration 1000 / 1500: loss 7.003221
iteration 1100 / 1500: loss 6.624018
iteration 1200 / 1500: loss 6.332740
iteration 1300 / 1500: loss 5.313708
iteration 1400 / 1500: loss 5.623894
That took 9.446967s
```

```
In [14]: 1 # A useful debugging strategy is to plot the loss as a function of
2 # iteration number:
3 plt.plot(loss_hist)
4 plt.xlabel('Iteration number')
5 plt.ylabel('Loss value')
6 plt.show()
```



```
In [15]: 1 # Write the LinearSVM.predict function and evaluate the performance on both the
2 # training and validation set
3 y_train_pred = svm.predict(X_train)
4 print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
5 y_val_pred = svm.predict(X_val)
6 print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

```
training accuracy: 0.377327
validation accuracy: 0.378000
```


In [16]:

```
1 # Use the validation set to tune hyperparameters (regularization strength and
2 # learning rate). You should experiment with different ranges for the learning
3 # rates and regularization strengths; if you are careful you should be able to
4 # get a classification accuracy of about 0.4 on the validation set.
5 learning_rates = [1e-7, 5e-5]
6 regularization_strengths = [2.5e4, 5e4]
7
8 # results is dictionary mapping tuples of the form
9 # (learning_rate, regularization_strength) to tuples of the form
10 # (training_accuracy, validation_accuracy). The accuracy is simply the fraction
11 # of data points that are correctly classified.
12 results = {}
13 best_val = -1 # The highest validation accuracy that we have seen so far.
14 best_svm = None # The LinearSVM object that achieved the highest validation rate.
15
16 #####
17 # TODO:
18 # Write code that chooses the best hyperparameters by tuning on the validation #
19 # set. For each combination of hyperparameters, train a linear SVM on the #
20 # training set, compute its accuracy on the training and validation sets, and #
21 # store these numbers in the results dictionary. In addition, store the best #
22 # validation accuracy in best_val and the LinearSVM object that achieves this #
23 # accuracy in best_svm.
24 #
25 # Hint: You should use a small value for num_iters as you develop your #
26 # validation code so that the SVMs don't take much time to train; once you are #
27 # confident that your validation code works, you should rerun the validation #
28 # code with a larger value for num_iters.
29 #####
30 # Your code
31 num_iters = 800 # test the para
32 for lr in learning_rates:
33     for rs in regularization_strengths:
34         svm = LinearSVM()
35         svm.train(X_train, y_train, learning_rate=lr, reg=rs, num_iters=num_iters)
36
37         y_train_pred = svm.predict(X_train)
38         train_acc = np.mean(y_train == y_train_pred)
39         y_val_pred = svm.predict(X_val)
40         vali_acc = np.mean(y_val == y_val_pred)
41
42         results[(lr, rs)] = (train_acc, vali_acc)
43
44         if best_val < vali_acc:
45             best_val = vali_acc
46             best_svm = svm
47
48 #####
49 # END OF YOUR CODE
50 #####
51
52 # Print out results.
53 for lr, reg in sorted(results):
54     train_accuracy, val_accuracy = results[(lr, reg)]
55     print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
56         lr, reg, train_accuracy, val_accuracy))
57
58 print('best validation accuracy achieved during cross-validation: %f' % best_val)
```

```
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.368061 val accuracy: 0.380000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.369980 val accuracy: 0.368000
```

```

lr 5.000000e-05 reg 2.500000e+04 train accuracy: 0.167918 val accuracy: 0.167000
lr 5.000000e-05 reg 5.000000e+04 train accuracy: 0.049531 val accuracy: 0.049000
best validation accuracy achieved during cross-validation: 0.380000

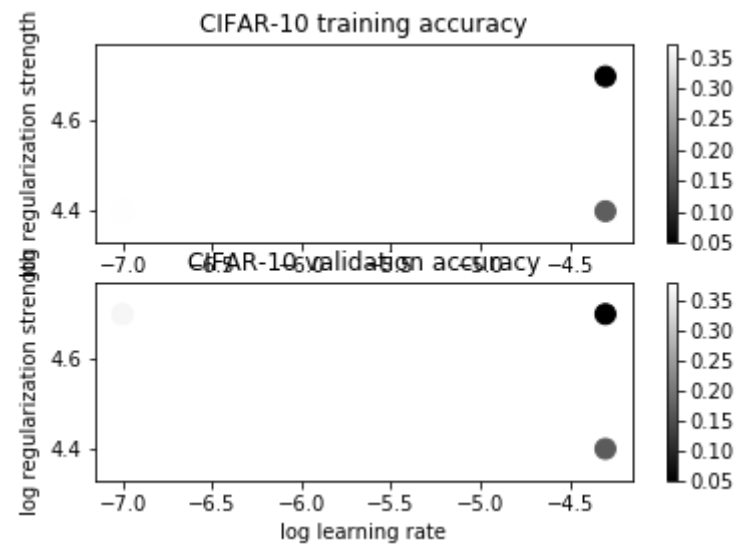
```

In [17]:

```

1 # Visualize the cross-validation results
2 import math
3 x_scatter = [math.log10(x[0]) for x in results]
4 y_scatter = [math.log10(x[1]) for x in results]
5
6 # plot training accuracy
7 marker_size = 100
8 colors = [results[x][0] for x in results]
9 plt.subplot(2, 1, 1)
10 plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
11 plt.colorbar()
12 plt.xlabel('log learning rate')
13 plt.ylabel('log regularization strength')
14 plt.title('CIFAR-10 training accuracy')
15
16 # plot validation accuracy
17 colors = [results[x][1] for x in results] # default size of markers is 20
18 plt.subplot(2, 1, 2)
19 plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
20 plt.colorbar()
21 plt.xlabel('log learning rate')
22 plt.ylabel('log regularization strength')
23 plt.title('CIFAR-10 validation accuracy')
24 plt.show()

```



In [18]:

```

1 # Evaluate the best svm on test set
2 y_test_pred = best_svm.predict(X_test)
3 test_accuracy = np.mean(y_test == y_test_pred)
4 print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)

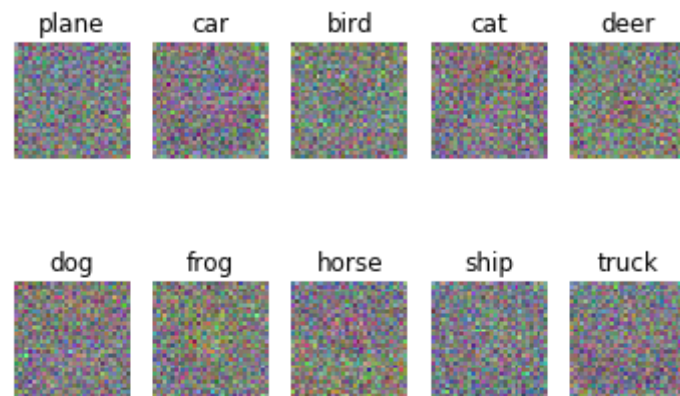
```

```
linear SVM on raw pixels final test set accuracy: 0.362000
```

```

In [19]: 1 # Visualize the learned weights for each class.
2 # Depending on your choice of learning rate and regularization strength, these may
3 # or may not be nice to look at.
4 w = best_svm.W[:-1,:] # strip out the bias
5 w = w.reshape(32, 32, 3, 10)
6 w_min, w_max = np.min(w), np.max(w)
7 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
8 for i in range(10):
9     plt.subplot(2, 5, i + 1)
10
11     # Rescale the weights to be between 0 and 255
12     wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
13     plt.imshow(wimg.astype('uint8'))
14     plt.axis('off')
15     plt.title(classes[i])

```



Inline question 2:

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way that they do.

Your answer: *fill this in*